



■ CS177 Bioinformatics: Software Development and Applications

Lecture 5: Introduction to Machine Learning and Data Science

Jie Zheng (郑杰)

PhD, Associate Professor

School of Information Science and Technology (SIST), ShanghaiTech University

Spring, 2024



Learning objectives



- After taking this lecture, you should be able to:
 - Define what machine learning is
 - Describe paradigms of machine learning
 - Explain when and why use machine learning
 - Explain how to measure the performance of machine learning methods
 - Tell what are the main challenges of machine learning
 - Describe how to train a machine learning model using regression as an example
 - Define what data science is
 - Describe the life cycle of a data science project





Introduction to Machine Learning






Example: Spam filtering



- **Input:** Incoming emails



- **Output:** "SPAM" or "NOT SPAM"
- A **binary classification** problem, i.e. having only 2 possible outputs






Traditional rule-based approaches



- A traditional method for spam filtering:
 - Study what spam typically looks like
 - A spam email usually contains some special key words in the subject, e.g. "4U" , "credit card" , ... (encoded as rules)
 - Write and test a spam detection algorithm





Issues with rule-based approaches



- Very labor intensive to build a system
- End up with a long list of rules, which is hard to maintain
- Rules may work very well for specific areas they cover, but cannot naturally handle uncertainty







What is machine learning?

- **Machine learning** is a “field of study that gives computers the ability to learn without being explicitly programmed” (Arthur Samuel, 1959)
- **Program by example**: Rather than program the computer to solve a task directly, with machine learning the computer will come up with its own program based on examples given to it

Traditional Programming



Machine Learning





Machine learning approach



- A machine “learns” on its own **without manually designed rules**
- We provide “training data” , i.e. many examples of (input x , output y), e.g.
 - A set of emails, and whether or not each is SPAM
 - A set of images, and whether or not each has a bird





What is machine learning? (continued)



- **Engineering-oriented view**: A computer program is said to **learn from experience E with respect to some task T and some performance measure P** , if its performance on T , as measured by P , improves with experience E (Tom Mitchell, 1997)
- A common theme is to solve a **prediction problem**:
 - Given an **input** x
 - Predict an “appropriate” **output** y
- Let us start with a canonical example





Example: Spam filtering (continued)




- **Task T:** To flag spam for new emails
- **Experience E:** The **training data**
 - Examples of spam emails (e.g. flagged by users)
 - Examples of regular (nospam) emails
- **Performance measure P:** To be defined
 - e.g. **accuracy** (ratio of correctly classified emails)



Machine learning algorithms



- A machine learning algorithm:
 - **Input:** Training data
 - “Learns” from the training data
 - **Output:** A prediction function that produces output y given x




The prediction function



- A **prediction function** takes input x and produces an output y
- We look for prediction functions that solve particular problems
- Machine learning techniques aim to find the **best** prediction function





Machine learning as a subfield of AI




Figure credit: <https://www.edureka.co/blog/ai-vs-machine-learning-vs-deep-learning/>





Paradigms of machine learning



- **Supervised learning**: Given a set of feature/label pairs, find a rule (or function) that predicts the label associated with a previously unseen input
 - Examples: classification, regression
 - **Semisupervised learning**: Training dataset is partially labeled
- **Unsupervised learning**: Given a set of feature vectors (without labels), group them into “natural clusters” or create labels for the groups
 - Example: clustering
- **Reinforcement learning**: An area of machine learning concerned with how an agent ought to take actions in an environment so as to maximize some notion of cumulative reward. The algorithm (or agent) learns by interacting with its environment.






Performance Measures




Classification vs. regression

- **Classification** is a typical supervised learning task, e.g. the spam filter is trained with many example emails along with their class (spam or not spam) and it must learn how to classify new emails
- **Regression** is a task to predict a target numeric value, given a set of **features** called **predictor**, e.g. to predict the price of a car given its mileage, age, brand, etc.



Spam classification



Regression

How to measure the performance of a classification model (i.e. classifier)?



■ Training/test splitting

- **Training set**: only for **training** prediction functions
- **Test set**: only for **evaluating model performance**, and must be **independent** of training data
- Training/test set is usually **randomly** split
- But if the data contain **time series**, two sets should be created by splitting in time
 - Training set is everything before time T
 - Test set is everything after time T






Cross-validation



- An intuitive method to test the model when training:
 - Randomly make *training / validation* split on the original training data
 - Perform the machine learning algorithm using the *training* set
 - Estimate the “future” result with the *validation* set
- A very **simple** evaluation method 😊
- We might **waste too much** training data, and both too large or too small datasets will cause the model evaluation imprecise 😞





Cross-validation



- Make full use of a dataset with a **k -fold cross-validation**:
 1. uniformly divide data into k blocks
 2. for $j = 1$ to k
 3. train on blocks except for the j th block
 4. validate on the j th block
 5. evaluate the result
 6. average the results

original

1
2
3
4
5

2
3
4
5

1
3
4
5

1
2
4
5

1
2
3
5

1
2
3
4

training

1

2


3

4

5

test






Confusion matrix




- To evaluate a classifier, **accuracy** (i.e. ratio of correct predictions) may be biased for **skewed** datasets (e.g. when some classes are much more frequent than others)
- A better metric for classifier performance is to count the number of instances of class A classified as class B, e.g. for a **binary** classifier:



TP: true positive
TN: true negative
FP: false positive
FN: false negative

Identify digit 5 from the MNIST database of handwritten digit images (Fig. 3-2 of A. Géron book)







Precision/Recall tradeoff



$$\text{precision} = \frac{TP}{TP + FP}$$

$$\text{recall} = \frac{TP}{TP + FN}$$

- Adjusting the **decision threshold** (i.e. if score is above such a threshold it is assigned positive, or else to negative) would change relative values on the **precision/recall (PR) curve**






ROC (receiver operating characteristic) curve



- Common used for binary classifiers: TP rate (recall) vs FP rate
- Ideal classifiers have ROC AUC (area under the curve) equal to 1



$$\text{TPR} = \text{TP} / (\text{TP} + \text{FN})$$

$$\text{FPR} = \text{FP} / (\text{FP} + \text{TN})$$





■ Hyperparameters (or “Tuning Parameters”)

- Almost every learning algorithm has
 - At least one “hyperparameter” or “tuning parameter”
- You must tune these values
- Hyperparameters control various things
 - Model complexity (e.g. polynomial order)
 - Type of model complexity control (e.g. L1 vs L2 regularization)
 - Optimization algorithm (e.g. learning rate)
 - Model type (e.g. loss function, kernel type, ...)





Strengths and Challenges of Machine Learning



Why machine learning?



- Some problems are too difficult to program by hand:
 - Fuzzy and “noisy” patterns
 - Large amounts of data
 - Mechanisms are still unclear
 - Lack general theories to guide programming
- Biomedical sciences (such as drug discovery) are often faced with the above challenges
- Other applications:
 - Face recognition
 - Spam filtering
 - Medical diagnosis
 - Natural language processing
 - Product recommendation
 - Fraud detection
 - Weather forecasting





Strengths of machine learning



- Usually simplify code and perform better for complex problems
- Machine learning systems can adapt to new data
- Can help extract patterns and insights from large amounts of data, i.e. automating knowledge discovery





Main challenges: Dependency on data



- Data might be insufficient, noisy or biased
 - Garbage in, garbage out
 - **Insufficient quantity of training data**: More data would provide more complete features learned by the model, and thus generate more accurate results
 - **Nonrepresentative training data**: Representative training data could generalize the learning algorithms to the expected new cases
 - e.g. sampling bias due to “unfair” sampling methods
 - **Poor-quality data**: Errors, outliers, missing values or noise in the data will make it harder for the system to detect underlying patterns



Generate a picture of Elon Musk




Sure



From Google' s Gemini AI







Main challenges: ML is often black box



- A model is often a “black box” : More **interpretability** is needed
 - Why did you do that?
 - Why did you not do that?
 - When do you succeed or fail?
 - How do I correct an error?




Overfitting

- **Overfitting the training data:**

- Performance is good on the training data
- But it does not generalize well


- Potential reasons:

- **High dimensionality** (too many features)
- **Model is too complex** (too many parameters)
- Noise in the training data



Underfitting

- A model **underfits the training data** when it is too simple to learn the underlying structure of the data
- The model
 - cannot fit most of the data, and has poor training performance
 - should be made more powerful to catch richer features



A model with only one parameter (hair length): Sun Y.S. is a boy.




Training Regression Models



■ Review of regression



- Given a set of attributes $x : (x_1, x_2, x_3, \dots, x_n)$ of an object, estimate the mapping function from input x to a continuous output variable y base on training examples.





What is linear regression?



- **Linear Regression model** : A linear model makes a prediction by simply computing a **weighted sum** of the input features x :

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \theta_n x_n$$

- $\hat{y}$: predicted value
- n : number of the features
- x_i : the i^{th} feature value
- θ_j : the j^{th} model parameter (include bias term θ_0)





Linear regression



- A linear regression model (in **vectorized** form):

$$\hat{y} = \theta \cdot x = \theta^T x$$

- θ : model parameter vector, containing a bias term and feature weights from θ_0 to θ_n
- x : feature vector of an instance (with $x_0 = 1$)
- $\theta \cdot x$ is the dot product of the vectors θ and x which is equal to : $\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \theta_n x_n$



Linear regression's objective

- How to obtain a proper linear regression model from training examples?
 - Training a model means setting its parameters so that **the model best fits the training dataset**
- What does “best fits” mean?
 - The most common performance measure of a regression model is the **Root Mean Square Error (RMSE)**:

$$\text{RMSE}(\mathbf{X}, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m \left(h(\mathbf{x}^{(i)}) - y^{(i)} \right)^2}$$

- In practice, it is simpler to minimize the **Mean Square Error (MSE)**:

$$\text{MSE}(\mathbf{X}, h_{\boldsymbol{\theta}}) = \frac{1}{m} \sum_{i=1}^m \left(\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)} \right)^2$$

- Lower RMSE or MSE scores represent better model fitting





Linear regression' s objective (continued)




- After combining the MSE and linear regression model, we can build a cost function :

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (\theta^T \mathbf{x}^{(i)} - y^{(i)})^2$$

- Our objective is to **minimize the MSE cost** by tuning the parameter θ , i.e. to make the model fit to the training samples.






Gradient Descent



- Goal: To decrease the MSE cost via updating parameter θ
- Approach: **Gradient Descent**
 - Filling θ with random values
 - Keep changing θ to reduce the objective function $J(\theta)$





Gradient Descent



- How to change the parameter θ ?

$$\theta_j^{k+1} = \theta_j^k - \alpha \frac{\partial}{\partial \theta_j} J(\boldsymbol{\theta}^k)$$
$$\frac{\partial}{\partial \theta_j} J(\boldsymbol{\theta}^k) = \frac{2}{m} \sum_{i=1}^m (\boldsymbol{\theta}^T \mathbf{x}^{(i)} - y^{(i)}) x_j^{(i)}$$

- k : iteration count
- θ_j^k : the j^{th} parameter in the k^{th} iteration
- α : step size (also known as **learning rate**)







Learning rate in Gradient Descent



- If the learning rate is too small, it is hard to converge






Learning rate in Gradient Descent



- If the learning rate is too large, results are unstable






Pitfalls of Gradient Descent



Some cost functions may cause problems:

- **Local Minimum** : It cannot reach the optimal solution by being stuck in local minimum
- **Plateau** : It takes a long time to cross the plateau





Other Gradient Descent methods





- **Batch Gradient Descent:** Instead of individually computing the partial derivatives, solve them all in one in vectorized form
- **Stochastic Gradient Descent (SGD):** Pick one random instance from training set at each iteration to compute gradient based on that instance
- **Mini-batch Gradient Descent:** Compute the gradients on small random subsets of instances



Polynomial Regression

- Linear Regression sometimes cannot fit the training sample well, if the data is nonlinear
- How to use a linear model to fit nonlinear data?
- **Polynomial Regression** is a technique that adds the powers of the original features as new features, and then train a linear model on this extended set of features




Polynomial Regression



- An example of Polynomial Regression:


$$\hat{y} = \theta_2 x_1^2 + \theta_1 x_1 + \theta_0$$

- The example input x only has one feature x_1
- The degree of this model is 2




Learning curves

- **Learning curves** are plots of the model performance (e.g. cost function RMSE) on the training and validation sets as the training set size (or training iteration) changes
 - e.g. training two models on the same data:




Linear regression: **Underfitting**



Polynomial regression: **Overfitting**

Regularized linear model

- Can we do better for point fitting?
 - Increase the degree of model
- High-degree Polynomial Regression:
 - Low error
 - **Overfitting!**
- **Regularization**
 - Constraining the model to make it simpler and harder for it to overfit the data
 - Example: to regularize a polynomial model by reducing the number of polynomial degrees





Classical machine learning methods



- Support vector machine (SVM)
- Decision tree
- Ensemble learning (e.g. Random Forest)
- Dimensionality reduction (e.g. PCA, t-SNE)
- Clustering algorithms (e.g. K-means, DBSCAN)
- Bayesian statistics
- Neural network





Overview of Data Science






Definition of Data Science




- Data science is a relatively new field, yet pervasive in all industries. With more data every day comes the challenge of managing and analyzing it
- **Data Science:**
 - means “the scientific study of the creation, validation and transformation of data to create meaning.”¹
 - is “the interdisciplinary field of inquiry in which quantitative and analytical approaches, processes, and systems are developed and used to extract knowledge and insights from increasingly large and/or complex sets of data.”²

1. Data Science Association. <http://www.datascienceassn.org/about-data-science> Accessed March 24, 2019
2. NIH Strategic Plan for Data Science. Bethesda, MD. National Institutes of Health. 2018. Accessed March 24, 2019






Related Fields



- This Venn diagram shows the important overlaps among multiple fields
- Clearly, data science is the blend between mathematics, statistics and computer science, coupled with domain knowledge






Data Science Career (in the USA)



- Due to the shortage of senior data scientists in the pipeline, **all data scientists are in high demand**, in all domains (including healthcare)
- Starting salaries are said to be between \$108 - \$131 K. Data scientist is called the “sexiest job of the 21st century” (Davenport, Patil, *Harvard Business Review*, 2012).
- Skills requested:



Data Science Life Cycle



- **Data Life Cycle** describes the various stages a dataset traverses as it undergoes scientific collection and investigation
 - Typically used to guide data management decisions and practices
- **Data Science Life Cycle** is an extension of “the Data Life Cycle”
 - Describe the complete process of a data science project
 - Focus on the generation of scientific findings, including computational components, inferential methodology, etc.




Figure: Example of a data life cycle. Source: Berman, F. et al. CACM, 2018.






The OSEMN Framework



- **1. Obtain data:** Receive data from various sources, e.g. databases, web APIs, Excel files
- **2. Scrub data:** Clean and filter the data, e.g. merge or split data, handle missing data, standardize format across all data
- **3. Explore data:** Inspect data and its properties; compute descriptive statistics to extract features and test significant variables; use data visualization to identify significant patterns and trends in the data; figure out business questions and transform them into a data science question
- **4. Model data:** Dimensionality reduction; forecast future values, classification, clustering
- **5. Interpret data:** Present results to users with actionable insights; visualize findings; tell a story



Originally by Hilary Mason and Chris Wiggins

Blog of Dr. Cher Han Lau (2019): <https://towardsdatascience.com/5-steps-of-a-data-science-project-lifecycle-26c50372b492>



Summary



- In this lecture we have learned:
 - What machine learning is
 - Types of machine learning systems, e.g. supervised and unsupervised learning, etc.
 - Methods to measure the performance of classifiers, e.g. cross-validation, confusion matrix, ROC curve, etc.
 - Why use machine learning: when traditional methods fail
 - Some challenges of machine learning
 - How to train a machine learning model, using linear regression as an example
 - What is data science and the life cycle of a data science project





References



- Aurélien Géron. Chapters 1, 3 and 4, “Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow” , O’ Reilly, 2019 (Ed. 2)
- Thomas H. Davenport and D.J. Patil (2012). Data Scientist: The Sexiest Job of the 21st Century. Harvard Business Review.
- Victoria Stodden (2020). The Data Science Life Cycle: A Disciplined Approach to Advancing Data Science as a Science. Communications of The ACM, Vol. 63, No. 7, pp. 58-66.





Appendix: Machine Learning in Bioinformatics






Machine learning techniques



- **Hidden Markov model (HMM):** sequence and profile alignment
- **Support vector machine (SVM):** protein fold recognition, mutation prediction
- **Bayesian networks:** inference of gene regulatory network from gene expression data
- **Neural networks:** secondary structure prediction, fold prediction, contact prediction)
- **Deep learning:** Autoencoders for single-cell RNA-seq data processing and data analysis, AlphaFold 2 for protein structure prediction






Clustering of microarray data




Time:	Time X	Time Y	Time Z
Gene 1	10	8	10
Gene 2	10	0	9
Gene 3	4	8.6	3
Gene 4	7	8	3
Gene 5	1	2	3




- **Input data:** **intensity matrix** (expression levels of genes at measured time points), which allows biologists to
 - Make correlations between different genes
 - Understand how genes might be related in functions
- **Application:** inferring gene function
 - We estimate the distance between every two genes
 - Genes with a small distance have similar expression patterns, and thus might be functionally related
 - Clustering reveals groups of functionally related genes





Classification of molecular interactions



Zheng et al. (2020) "Predicting drug-protein interaction using quasi-visual question answering system." *Nature Machine Intelligence*, Vol. 2, pp. 134-140

An end-to-end deep learning framework to predict the interactions (as a binary classification task), by representing:

- Proteins with a two-dimensional distance map from monomer structures (image)
- Drugs with molecular linear notation (string)

