



■ CS177 Bioinformatics: Software Development and Applications

Lecture 14: Genomics

Jie Zheng (郑杰)

PhD, Associate Professor

School of Information Science and Technology (SIST), ShanghaiTech University

Spring, 2024





Learning objectives



- At the end of this lecture, you will be able to:
 - Define genomics and related basic concepts (e.g. functional genomics)
 - Name a few next-generation sequencing (NGS) technologies
 - Describe and explore some open data resources of genomics (e.g. cancer genomics projects)
 - Explain main ideas of genome-wide association studies (GWAS)
 - Explain the key idea of using convolutional neural network for modeling protein-DNA binding specificities
 - Describe what DeepBind can do and how it works





Overview of Genomics


PPT slides of this module are courtesy from Dr. Lichun Jiang, SIAIS, ShanghaiTech







Information flow in cells: Central dogma of modern biology

DNA double helix



<https://media.buzzle.com/>



https://en.wikipedia.org/wiki/Central_dogma_of_molecular_biology

<https://commons.wikimedia.org/w/index.php?curid=80147407>



What are genome and genomics



- An organism's complete set of DNA is called its **genome**. Virtually every single cell in our body contains a complete copy of the approximately 3 billion DNA base pairs, or letters, that make up the human genome.
- **Genomics** is an interdisciplinary field of biology focusing on the structure, function, evolution, mapping, and editing of genomes.

<https://ghr.nlm.nih.gov/primer/hgp/genome>

<https://en.wikipedia.org/wiki/Genomics>






DNA sequencing: technology to read DNA



上海科技大学
ShanghaiTech University



<https://www.genome.gov/genetics-glossary/acgt>


立志成才 报國裕民

DNA sequencing technologies




上海科技大学
ShanghaiTech University


Sanger




PacBio



NGS



Nanopore



立志成才 报 国 裕 民



Next-generation sequencing (NGS)




Technology	Read length (bp)	Reads per run	Time per run	Cost per megabase	Accuracy
Roche 454	700	1 million	1 day	\$10	99.9%
Illumina	50-250	<3 billion	1-10 days	~\$0.10	98%
SOLiD	50	~1.4 billion	7-14 days	\$0.13	99.9%
Ion Torrent	200	<5 million	2 hours	\$1	98%
Pacific Biosciences	2900	<75,000	<2 hours	\$2	99%
Sanger	400-900	N/A	<3 hours	\$2400	99.9%




Randomly fragment genomic DNA

Library preparation




Samples immobilized on
surface of a flow cell (8 lanes)

Solid phase amplification



- Bridge amplification (inverted U) generates clusters on surface of flow cell
- ~Ten million single-molecule clusters per square centimeter

Sequencing by synthesis



- Each cycle: add polymerase, one labeled deoxynucleoside triphosphate (dNTP) at a time (four labeled dNTPs per cycle)
- Image fluorescent dyes
- Call nucleotide
- Enzymatic cleavage to remove

■ Sequencing by Illumina


Disadvantage:

- Short read length (~150 bases)


Advantages:

- Very fast
- Low cost per base
- Large throughput, up to 1 gigabase/experiment
- Short read length makes it appropriate for resequencing
- No need for gel electrophoresis
- High accuracy
- All four bases are present at each cycle, with sequential addition of dNTPs. This allows homopolymers to be accurately read

https://www.bilibili.com/video/BV1v34y147nj/?spm_id_from=333.788.recommend_more_video.3



Third-generation sequencing (TGS)



Goodwin, S., McPherson, J. D., & McCombie, W. R. (2016). Coming of age: Ten years of next-generation sequencing technologies. *Nature Reviews Genetics*, 17(6), 333–351.






Third-generation sequencing (TGS)



- TGS is also known as long-read sequencing
 - In second-generation sequencing (SGS), DNA is fragmented into tiny bits, amplified and then sequence
 - TGS methods do not break or amplify DNA; instead, they sequence a single DNA molecule directly
- Strengths:
 - Offer capability of single molecule real-time sequencing of **longer** reads
 - Detect DNA modifications without any assay
 - Higher speed of sequencing than SGS, esp. in metagenomics
 - Better performance for transcript detection, mRNA abundance estimation, etc.
- Main challenge:
 - **High error rates**, due to instability of the molecular machinery involved
- Representative TGS technologies:
 - PacBio (Pacific Biosciences) SMRT (single molecule real-time) sequencing
 - Oxford nanopore sequencing





Third-generation sequencing (TGS)



上海科技大学
ShanghaiTech University

Aa Pacific Biosciences


SMRTbell template

Two hairpin adapters allow continuous circular sequencing




ZMW wells

Sites where sequencing takes place



Labelled nucleotides

All four dNTPs are labelled and available for incorporation



Modified polymerase

As a nucleotide is incorporated by the polymerase, a camera records the emitted light

PacBio output


A camera records the changing colours from all ZMWs; each colour change corresponds to one base



Ab Oxford Nanopore Technologies

Leader-Hairpin template

The leader sequence interacts with the pore and a motor protein to direct DNA, a hairpin allows for bidirectional sequencing




Motor protein

Alpha-hemolysin

A large biological pore capable of sensing DNA

Current

Passes through the pore and is modulated as DNA passes through




ONT output (squiggles)

Each current shift as DNA translocates through the pore corresponds to a particular k-mer


Sara Goodwin et al. (2016) Coming of age: ten years of next-generation sequencing technologies, Nature Reviews Genetics.

立志成才报 13 国裕民


Cost of DNA sequencing dropped more than Moore's Law in the past decade



GenBank and WGS Statistics



Production cost per 30x Human genome over 18 years



<https://www.ncbi.nlm.nih.gov/genbank/statistics/>

<http://enseqlopedia.com/>



Data Resources of Genomics

PPT slides of this module are courtesy from Dr. Lichun Jiang, SIAIS, ShanghaiTech



It took 20 years from a human genome draft to a complete one



2001



2022

- International Human Genome Consortium. Initial sequencing and analysis of the human genome. Nature, 409: 860-921. 2001.
- Venter, J.C. Adams, M.D., Myers, E.W., Li, P.W., Mural, R.J., et al. The sequence of the human genome. Science, 291: 1304-1351. 2001.
- The complete sequence of a human genome. Science. 2022 March 376:44-53




```
[jianglc@SIAIS-mgt gene_reference]$ cat test.fasta
>Random
GTGGGGTAGGGAGGTGCCTCACACATTTACGGAGCTTGATCTCATCCACAGCTGTCTCCGTGTAATGCCCGCACTCTTCACCACTTTGAGGGCCACAAAGCGCTTGCGCCTGCAA
CCCAGGGACAGGATCCCAGCCAGTGCTGGGGCCCCGGAGGCGAAGGAGGCACTCACTGGATGTCCCAGCAGAGCCAGACGGTGGAGAAGTGGCCCCAGCCCAGTTTGCGCACCACGT
```

[illegible][illegible]

```
(CRISPRmatch) [jianglc@SIAIS-mgt alignment]$ cat ../PeakCall/JLCH_13_S5 |head -n4
```

chr1	9958	10500	peak_0	1000	.	668.096802	4.995196	-1	82
chr1	28915	29465	peak_1	991	.	544.855103	4.598983	-1	423
chr1	32457	33073	peak_2	1000	.	922.261353	5.002809	-1	232

```
[jianglc@STAIS-mgt Quantification]$ head -n3 expression.txt
```


Gene_ID	Reads_T1	Reads_T2	Reads_T3	Reads_C1	Reads_C2	Reads_C3
ENSRNOG000000047964	0	2	20	108	201	189
ENSRNOG000000050370	11	13	20	900	1003	980



Functional genomics



Functional genomics is the study of how genes and intergenic regions of the genome contribute to different biological processes. (<https://www.ebi.ac.uk/training/online/courses/functional-genomics-i-introduction-and-design/what-is-functional-genomics/>)




Based on an image by Darryl Leja (NHGRI), Ian Dunham (EBI), Michael Pazin (NHGRI)

<https://www.encodeproject.org/>



**Bioinformatics tools
evolve with biotech!**

- **Sequence**
- **Quantification**
- **Interaction**
- **Association and causality**






ENCODE on model organisms



Abascal, F., Acosta, R., Addleman, N.J. *et al.* *Nature* 2020; 583:693–698.



CCLE on cancer cell lines



上海科技大学
ShanghaiTech University

[←](#) [→](#) [↻](#) [depmap.org/portal/cell_line/ACH-000364?tab=mutation](#)

应用

统一身份认证

学术

学校和所资源

statistics

database

BIOINFORMATICS

NGS experiments

researchers

vendors for device

animals

E-supply

bioinformatics

其他杂

 **depmap portal**

Tools ▾

Downloads

About ▾

Questions? Feedback?

U2OS

Depmap ID: ACH-000364 CCLL Name: U2OS_BONE COSMIC ID: 909776 Sanger ID: 1145 Cellosaurus RRID: CVCL_0042
Aliases: U-2 OS, U2-OS, U-2 OS

Disease

Bone Cancer

Disease Subtype

Osteosarcoma

Source

ATCC

Gender

Female

Lineage

Bone

Lineage Subtype

Osteosarcoma

Primary/Metastasis

Primary

U2OS sequenced and profiled in following datasets

Loss-of-function screens

CRISPR (Avana) Public 20Q3
Combined RNAi (Broad, Novartis, Marcotte)
RNAi (Broad)
RNAi (Novartis)

Omics Characterizations

Expression Public 20Q3
Copy Number Public 20Q3
Mutation Public 20Q3
Fusion 20Q3
Translocation
Protein Array
Methylation (1kb upstream TSS)
Copy Number (Absolute)
Proteomics

Compound Viability screens

Drug sensitivity AUC (Sanger GDSC1)
Drug sensitivity AUC (Sanger GDSC2)
Drug sensitivity IC50 (Sanger GDSC1)
Drug sensitivity IC50 (Sanger GDSC2)
Drug sensitivity AUC (CTD^2)
Drug sensitivity (PRISM Repurposing Primary Screen) 19Q4
Drug sensitivity AUC (PRISM Repurposing Secondary Screen) 19Q4
Drug sensitivity replicate-level dose (CTD^2)
Drug sensitivity replicate-level dose (Sanger GDSC1)
Drug sensitivity replicate-level dose (Sanger GDSC2)
Drug sensitivity dose-level (PRISM Repurposing Secondary Screen) 19Q4
Drug sensitivity replicate-level dose (PRISM Repurposing Secondary Screen) 19Q4

Top 10 preferentially essential genes

CRISPR (Avana) Public 20Q3

Combined RNAi (Broad, Novartis, Marcotte)



<https://depmap.org/portal/>

立志成才 报国裕民

Cancer genomics project



上海科技大学
ShanghaiTech University



<http://compbio.med.harvard.edu/chromothripsis/>

<https://cancer.sanger.ac.uk/cosmic>

The ICGC/TCGA Pan-Cancer Analysis of Whole Genomes Consortium.
[Nature](#). 2020; 578(7793): 82–93. Pan-cancer analysis of whole genomes


立志成才 报 国 裕 民



Human cell atlas



上海科技大学
ShanghaiTech University



**HUMAN
CELL
ATLAS**

A map of human cells

Mission:

To create comprehensive reference maps of all human cells—the fundamental units of life—as a basis for both understanding human health and diagnosing, monitoring, and treating disease

<https://www.humancellatlas.org/>





立志成才报²¹国裕民



Metagenomics



上海科技大学
ShanghaiTech University



Fan Wu et al. Nature. 2020; 579(7798): 265–269

<https://idseq.net/>

<http://merenlab.org/2020/07/27/history-of-metagenomics/>

Cell (2019) 176(3):649-662




Genome-Wide Association Studies (GWAS)




Quantitative Trait Loci (QTL) Mapping



- Association between
 - Genetic polymorphism (SNPs)
 - Phenotypic variation (e.g. height)



Human GWAS study



- Largest genome-wide association study (563,085) of blood cell traits to date
- Empiric assessments of omnigenic and infinitesimal models of polygenic variation
- Functional insights into how genetic variants impact human hematopoiesis
- Assessment of the effect of polygenic trait scores upon blood diseases

Dragana Vuckovic et al. Cell (2020), 182(5):1214-1231

- **Genome-wide association studies (GWAS)** uses SNP markers to identify disease loci
 - In **family-based** designs, markers are measured in a proband (i.e. the affected individual) and unaffected individuals to identify the differences in the frequency of variants
 - In **population-based** designs, a large number of **unrelated** cases and controls are studied (typically hundreds or thousands in each group). Larger sample sizes increase statistical power
- It is difficult to identify the genetic causes of common human diseases that involve **multiple** genes, each of which may make only a small contribution to the disease risk






Definition of GWAS




- GWAS is the association analysis performed with a panel of polymorphic markers (e.g. SNPs) adequately spaced to capture most of the linkage disequilibrium (LD) information in the entire genome in the population under study. Usually, a GWAS involves 500k – 4m SNPs



Test each SNP for
association with disease
phenotype



Clinical translation of findings from GWAS




McCarthy MI et al. Genome-wide association studies for complex traits: consensus, uncertainty and challenges. Nature Review Genetics, 2008 May; 9(5):356-69.



Manhattan plots



- Manhattan plots display GWAS findings with respect to the genomic locations of SNPs, highlighting signals of particular interest.
- The following is a Manhattan plot for type 2 diabetes






Software case study: PLINK



- PLINK is a free, open-source software toolset for whole-genome association analysis, designed to perform a range of basic, large-scale analyses in a computationally efficient manner
- Available at:
 - <http://zzz.bwh.harvard.edu/plink/> (old version)
 - <https://www.cog-genomics.org/plink/2.0/> (version 2.0)
- PLINK has numerous useful features for managing and analyzing genetic data
 - Data management
 - Summary statistics
 - Population stratification
 - Association analysis
 - Linkage disequilibrium and haplotype analysis
 - Shared segment analysis
 - Copy number analysis



Running PLINK with GUI



From Purcell et al. 2007






Case Study of Deep Learning in Genomics: DeepBind

PPT slides adapted from MIT course “Deep Learning in the Life Sciences”, Lecture 6 “Regulatory genomics”, by Prof. Manolis Kellis



Transcription factors (TFs) bind to specific DNA sequences to control gene expression






Enrichment-based motif discovery methods



Given a set of **co-regulated/functionally related** genes,
find common motifs in their promoter regions



- Align the promoters to each other using local alignment
- Use expert knowledge for what motifs should look like
- Find “median” string by enumeration (motif/sample driven)
- Start with conserved blocks in the upstream regions






Starting positions $\Leftrightarrow$ motif matrix





- Given aligned sequences, it is easy to compute a profile matrix
- Given a profile matrix, easy to find starting position probability





Deep convolutional neural network





Representing DNA sequence as 2D matrix




上海科技大学
ShanghaiTech University



Matrix representation of
DNA sequence
(darker = stronger)






Convolution – extracting invariant feature




Applying 4 bp sequence filter along the DNA matrix:




Yellow = high activity; blue = low activity



Convolution – extracting invariant feature (continued)



Rectification = ignore signals below some threshold.
Pooling = summary of each channel by max or average.




Prediction using extracted feature map



ChIP-seq, PBMs, SELEX Experiments DNA sequence

	A	T	G	C	A	G	C	A	N	N	N
A	■				■			■	■	■	■
T		■							■	■	■
G			■			■			■	■	■
C				■			■		■	■	■






DeepBind – overview of pipeline




上海科技大学
ShanghaiTech University


1. High-throughput experiments



2. Massively parallel deep learning




3. Community needs



Alipanahi et al. Nature Biotechnology 2015






DeepBind – details of inner workings



上海科技大学
ShanghaiTech University



Alipanahi et al. Nature Biotechnology 2015





Recap



- In this lecture we have learned:
 - Basic concepts and brief history of genomics
 - Technologies for DNA sequencing, especially the Next-Generation Sequencing (NGS)
 - Basic steps in a pipeline of GWAS
 - How to use PLINK software for GWAS
 - Key ideas of using CNN for DNA sequencing modeling
 - How DeepBind works to predict the binding specificities between proteins and DNA/RNA sequences





References



- Jonathan Pevsner (2015). Bioinformatics and Functional Genomics. (3rd Edition):
 - Chapter 9 "Analysis of Next-Generation Sequencing Data"
 - Chapter 20 "Human Genome"
 - Chapter 21 "Human Disease"
- Purcell, S. et al. PLINK: A tool set for whole-genome association and population-based linkage analyses. American Journal of Human Genetics, Vol. 81, pp. 559-575 (2007).
- Shendure, J. et al. DNA sequencing at 40 past, present and future. Nature 550(7676):345-353 (2017).
- Sara Goodwin et al. (2016) Coming of age: Ten years of next-generation sequencing technologies, Nature Reviews Genetics, 17(6), 333-351.
- Sapoval, N., Aghazadeh, A., Nute, M.G. et al. Current progress and open challenges for applying deep learning across the biosciences. Nature Communications **13**, 1728 (2022).
- Zou, J. et al. A primer on deep learning in genomics. Nature Genetics **51**, 12-18 (2019).
- Alipanahi, B. et al. Predicting the sequence specificities of DNA- and RNA-binding proteins by deep learning. Nature Biotechnology, **33**, 831-838 (2015). (**DeepBind**)





Appendix: Statistical Modeling in GWAS






Case-control study design



- Cases:
 - Ascertained for the phenotype of interest
 - Presumed to have a high prevalence of susceptibility alleles
- Controls:
 - Not ascertained for the phenotype
 - Presumed to have lower prevalence of such susceptibility alleles
- Misclassification bias: e.g. failure to identify **latent diagnoses** in the control group. It leads to loss of power, especially for common traits (e.g. high blood pressure)





Analysis outline



- Data cleaning / QC (quality control)
- Association test between genotype and phenotype:
 - Analyze each SNP separately
 - Use logistic or linear regression with additive models, including important covariates
 - Calculate a p-value for each SNP
- Control for population substructures, e.g. add principal components as covariates
- Visualization of results
 - Q-Q plots
 - Manhattan plots
- Control for multiple testing issue, e.g. Bonferroni correct, false discovery rate (FDR), Q-value





Odds ratios



- A measure of effect size, or strength of association
- Odds = $P / (1 - P)$, where P stands for probability of an event
- **Odds ratio** is the ratio of odds for an event (having a disease) under two comparing conditions (e.g. different genotypes at a SNP)
- Example (likelihood of a disease):
 - $P(D \mid \text{genotype "AT"}) = 0.8$
 - $P(D \mid \text{genotype "TT"}) = 0.2$
 - The odds ratio for getting the disease with genotype AT compared to TT is
$$OR = (0.8/0.2) / (0.2/0.8) = 16$$
 - The OR for AT individuals relative to an average population risk of 25% is:
$$OR = (0.8/0.2) / (0.25 / 0.75) = 12$$



- GWAS would often report ORs relative to the low-risk allele or lowest-risk genotype
- To turn this into a meaningful risk estimate, the **prevalence of the disease** and the **genotype frequencies** must be taken into account
- $P(D) = \text{prevalence}$
$$= P(D|AA)P(AA) + P(D|Aa)P(Aa) + P(D|aa)P(aa)$$
- For each SNP, an odds ratio and associated p-value are calculated (e.g. using the chi-square test)
- Significance of the association is measured by: $-\log_{10}(P\text{-value})$





Multiple hypothesis testing issue



- A conventional $\alpha = 0.05$ threshold for p-values assumes a 5% chance of false-positive finding due to chance
 - When performing only one test, this is reasonable
 - However, when performing many (e.g. 1,000,000) tests, which is common in GWAS, this will lead to many false positives
- Example: Consider a case with 20 hypotheses to test, and the threshold of significant p-value is $\alpha = 0.05$. The probability of observing at least one significant result is:

P(at least one significant result)

$$= 1 - P(\text{no significant results}) = 1 - (1 - 0.05)^{20} \approx 0.64$$

- So, with 20 tests being considered, there is a 64% chance of observing at least one “significant” result, even if none of the tests is actually significant
- Correction methods:
 - **Bonferroni correction**: sets the significance cut-off at α/n , where n is the number of tests
 - **False discovery rate (FDR)**: expected proportion of false positives among all results claimed as significant
 - **Q-value**: The minimum FDR at which the test may be called significant, serving as a p-value that has been adjusted for the FDR

